

CSC373

Week 2: Greedy Algorithms

Announcements

- First tutorial tomorrow!
- Details on Piazza
- First Assignment to be posted tomorrow (May 19) after tutorial
- Due **June 1**

Recap

- Divide & Conquer

- Master theorem
- Counting inversions in $O(n \log n)$
- Finding closest pair of points in \mathbb{R}^2 in $O(n \log n)$
- Fast integer multiplication in $O(n^{\log_2 3})$
- Fast matrix multiplication in $O(n^{\log_2 7})$
- Finding k^{th} smallest element (in particular, median) in $O(n)$

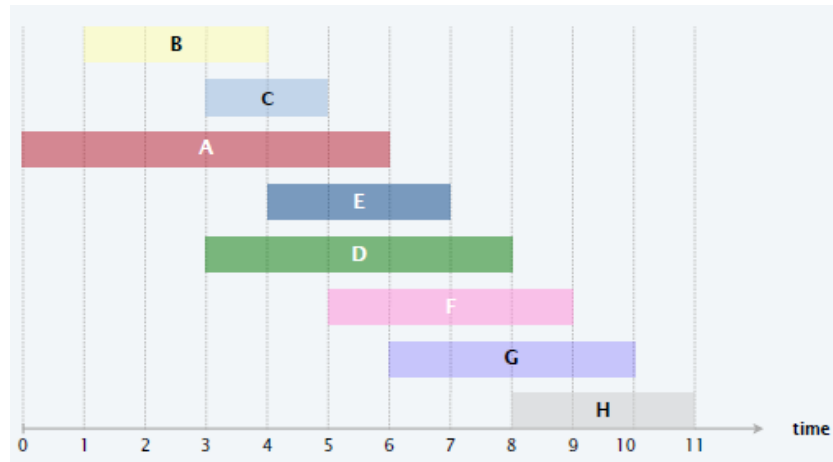
Greedy Algorithms

- Greedy/myopic algorithm outline
 - **Goal:** find a solution x maximizing/minimizing objective function f
 - **Challenge:** space of possible solutions x is too large
 - **Insight:** x is composed of several parts (e.g., x is a set or a sequence)
 - **Approach:** Instead of computing x directly...
 - Compute it one part at a time
 - Select the next part “greedily” to get the most immediate “benefit” (this needs to be defined carefully for each problem)
 - Polynomial running time is typically guaranteed
 - Need to prove that this will always return an optimal solution despite having no foresight

Interval Scheduling

- **Problem**

- Job j starts at time s_j and finishes at time f_j
- Two jobs i and j are compatible if $[s_i, f_i)$ and $[s_j, f_j)$ don't overlap
 - Note: we allow a job to start right when another finishes
- **Goal:** find maximum-size subset of mutually compatible jobs



Interval Scheduling

- Greedy template

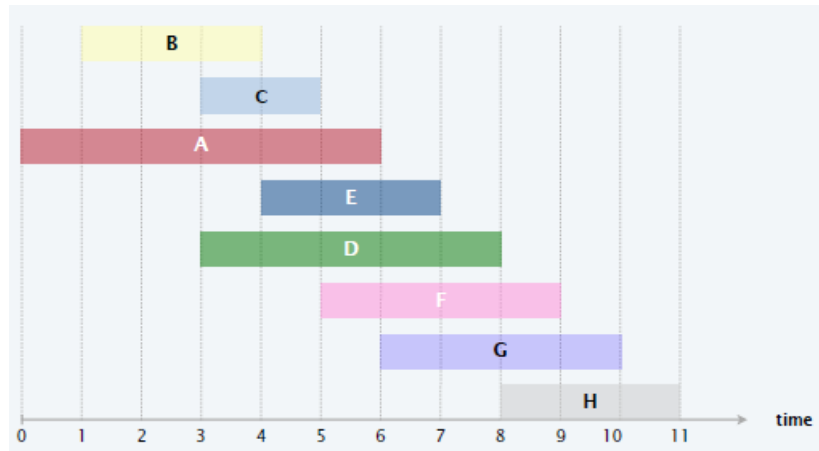
- Consider jobs in some “natural” order
- Take a job if it’s compatible with the ones already chosen

- What order?

- Earliest start time: ascending order of s_j
- Earliest finish time: ascending order of f_j
- Shortest interval: ascending order of $f_j - s_j$
- Fewest conflicts: ascending order of c_j , where c_j is the number of remaining jobs that conflict with j

Example

- **Earliest start time:** ascending order of s_j
- **Earliest finish time:** ascending order of f_j
- **Shortest interval:** ascending order of $f_j - s_j$
- **Fewest conflicts:** ascending order of c_j , where c_j is the number of remaining jobs that conflict with j



Interval Scheduling

- Does it work?



Counterexamples for

earliest start time



shortest interval



fewest conflicts


Interval Scheduling

- Implementing greedy with earliest finish time (EFT)
 - Sort jobs by finish time, say $f_1 \leq f_2 \leq \dots \leq f_n$
 - $O(n \log n)$
 - For each job j , we need to check if it's compatible with *all* previously added jobs
 - Naively, this can take $O(n)$ time per job j , so $O(n^2)$ total time
 - We only need to check if $s_j \geq f_{i^*}$, where i^* is the *last added job*
 - For any jobs i added before i^* , $f_i \leq f_{i^*}$
 - By keeping track of f_{i^*} , we can check job j in $O(1)$ time
 - Running time: $O(n \log n)$

Interval Scheduling

- **Proof of optimality by contradiction**

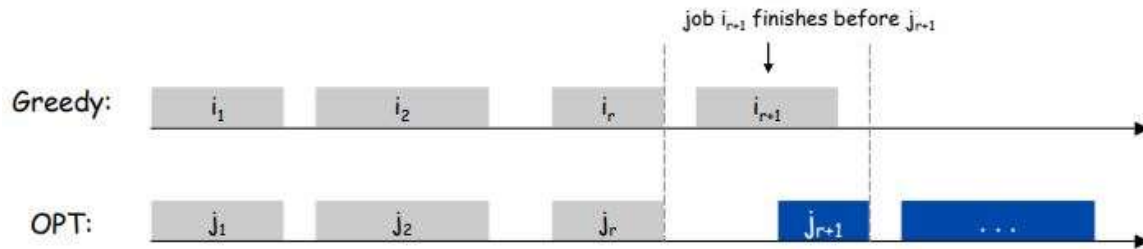
- Suppose for contradiction that greedy is not optimal
- Say greedy selects jobs i_1, i_2, \dots, i_k sorted by finish time
- Consider an optimal solution j_1, j_2, \dots, j_m (also sorted by finish time) which matches greedy for as many indices as possible
 - That is, we want $j_1 = i_1, \dots, j_r = i_r$ for the greatest possible r
- Both i_{r+1} and j_{r+1} must be compatible with the previous selection ($i_1 = j_1, \dots, i_r = j_r$)



Interval Scheduling

- **Proof of optimality by contradiction**

- Consider a new solution $i_1, i_2, \dots, i_r, i_{r+1}, j_{r+2}, \dots, j_m$
 - We have replaced j_{r+1} by i_{r+1} in our reference optimal solution
 - This is still feasible because $f_{i_{r+1}} \leq f_{j_{r+1}} \leq s_{j_t}$ for $t \geq r + 2$
 - This is still optimal because m jobs are selected
 - But it matches the greedy solution in $r + 1$ indices
 - This is the desired contradiction



Interval Scheduling

- **Proof of optimality by induction**

- Let S_j be the subset of jobs picked by greedy after considering the first j jobs in the increasing order of finish time
 - Define $S_0 = \emptyset$
- We call this partial solution **promising** if there is a way to extend it to an optimal solution by picking some subset of jobs $j + 1, \dots, n$
 - $\exists T \subseteq \{j + 1, \dots, n\}$ such that $O_j = S_j \cup T$ is optimal
- **Inductive claim:** For all $t \in \{0, 1, \dots, n\}$, S_t is promising
- If we prove this, then we are done!
 - For $t = n$, if S_n is promising, then it must be optimal (**Why?**)
 - We chose $t = 0$ as our base case since it is “trivial”

Interval Scheduling

- **Proof of optimality by induction**

- S_j is *promising* if $\exists T \subseteq \{j + 1, \dots, n\}$ such that $O_j = S_j \cup T$ is optimal
- **Inductive claim:** For all $t \in \{0, 1, \dots, n\}$, S_t is promising
- **Base case:** For $t = 0$, $S_0 = \emptyset$ is clearly promising
 - Any optimal solution extends it
- **Induction hypothesis:** Suppose the claim holds for $t = j - 1$ and optimal solution O_{j-1} extends S_{j-1}
- **Induction step:** At $t = j$, we have two possibilities:
 - 1) Greedy did not select job j , so $S_j = S_{j-1}$
 - Job j must conflict with some job in S_{j-1}
 - Since $S_{j-1} \subseteq O_{j-1}$, O_{j-1} also cannot include job j
 - $O_j = O_{j-1}$ also extends $S_j = S_{j-1}$

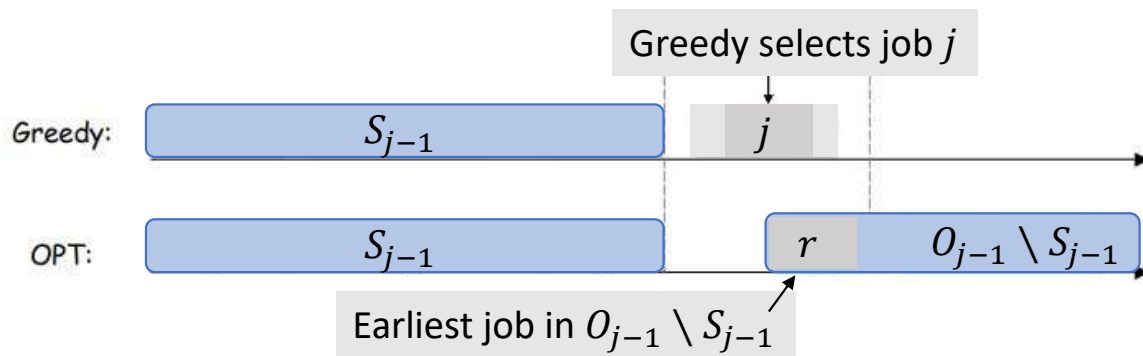
Interval Scheduling

- Proof of optimality by induction

➤ Induction step: At $t = j$, we have two possibilities:

2) Greedy selected job j , so $S_j = S_{j-1} \cup \{j\}$

- Consider the earliest job r in $O_{j-1} \setminus S_{j-1}$
- Consider O_j obtained by replacing r with j in O_{j-1}
- Prove that O_j is still feasible
- O_j extends S_j , as desired!



Contradiction vs Induction

- Both methods make the same claim
 - “The greedy solution after j iterations can be extended to an optimal solution, $\forall j$ ”
- They also use the same key argument
 - “If the greedy solution after j iterations can be extended to an optimal solution, then the greedy solution after $j + 1$ iterations can be extended to an optimal solution as well”
 - For proof by induction, this is the key induction step
 - For proof by contradiction, we take the greatest j for which the greedy solution can be extended to an optimal solution, and derive a contradiction by extending the greedy solution after $j + 1$ iterations

Interval Partitioning

- **Problem**

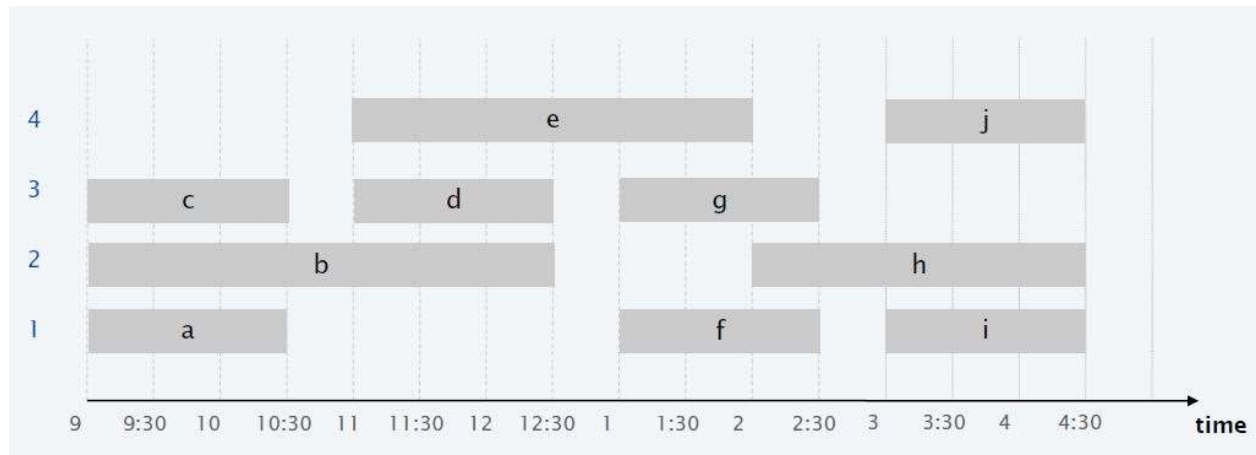
- Job j starts at time s_j and finishes at time f_j
- Two jobs are compatible if they don't overlap
- **Goal:** group jobs into fewest partitions such that jobs in the same partition are compatible

- **One idea**

- Find the maximum compatible set using the previous greedy EFT algorithm, call it one partition, recurse on the remaining jobs.
- Doesn't work (check by yourselves)

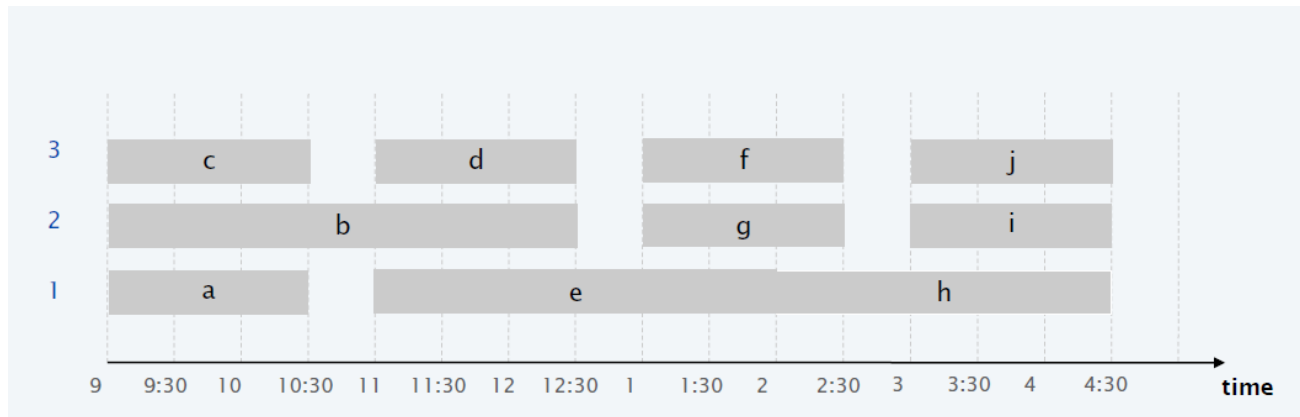
Interval Partitioning

- Think of scheduling lectures for various courses into as few classrooms as possible
- This schedule uses **4** classrooms for scheduling 10 lectures



Interval Partitioning

- Think of scheduling lectures for various courses into as few classrooms as possible
- This schedule uses **3** classrooms for scheduling 10 lectures



Interval Partitioning

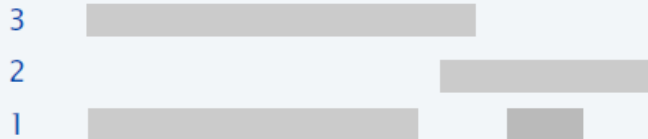
- Let's go back to the **greedy template**!
 - Go through lectures in some “natural” order
 - Assign each lecture to an (**arbitrary?**) compatible classroom, and create a new classroom if the lecture conflicts with every existing classroom
- **Order of lectures?**
 - **Earliest start time**: ascending order of s_j
 - **Earliest finish time**: ascending order of f_j
 - **Shortest interval**: ascending order of $f_j - s_j$
 - **Fewest conflicts**: ascending order of c_j , where c_j is the number of remaining jobs that conflict with j

Interval Partitioning

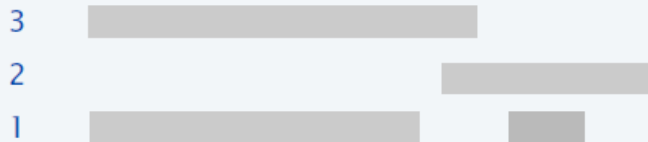
counterexample for earliest finish time



counterexample for shortest interval



counterexample for fewest conflicts



- At least when you assign each lecture to an arbitrary compatible classroom, three of these heuristics do not work.
- The fourth one works! (next slide)

Interval Partitioning

EARLIESTSTARTTIMEFIRST($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

SORT lectures by start time so that $s_1 \leq s_2 \leq \dots \leq s_n$.

$d \leftarrow 0$  number of allocated classrooms

FOR $j = 1$ **TO** n

IF lecture j is compatible with some classroom

 Schedule lecture j in any such classroom k .

ELSE

 Allocate a new classroom $d + 1$.

 Schedule lecture j in classroom $d + 1$.

$d \leftarrow d + 1$

RETURN schedule.

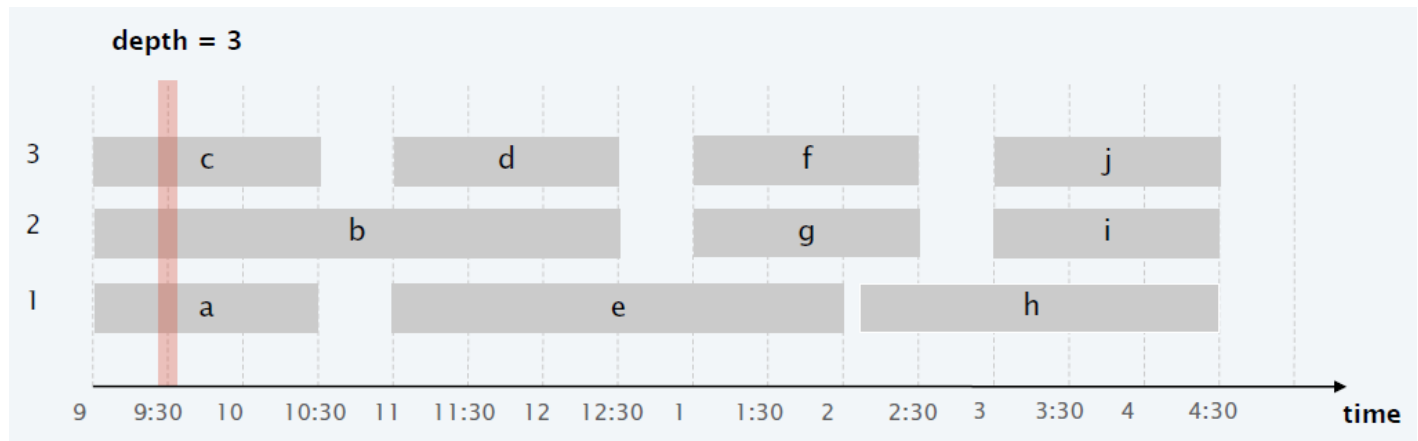
Interval Partitioning

- Running time

- **Key step:** check if the next lecture can be scheduled at some classroom
- Store classrooms in a priority queue
 - key = latest finish time of any lecture in the classroom
- Is lecture j compatible with some classroom?
 - Same as “Is s_j at least as large as the minimum key?”
 - If yes: add lecture j to classroom k with minimum key, and increase its key to f_j
 - Otherwise: create a new classroom, add lecture j , set key to f_j
- $O(n)$ priority queue operations, $O(n \log n)$ time

Interval Partitioning

- **Proof of optimality (lower bound)**
 - # classrooms needed \geq “depth”
 - depth = maximum number of lectures running at any time
 - Recall, as before, that job i runs in $[s_i, f_i)$
 - Claim: our greedy algorithm uses only these many classrooms!



Interval Partitioning

- **Proof of optimality (upper bound)**

- Let $d = \#$ classrooms used by greedy
- Classroom d was opened because there was a lecture j which was incompatible with some lectures already scheduled in each of $d - 1$ other classrooms
- All these d lectures end after s_j
- Since we sorted by start time, they all start at/before s_j
- So, at time s_j , we have d mutually overlapping lectures
- Hence, $\text{depth} \geq d = \#$ classrooms used by greedy ■
- Note: before we proved that $\#$ classrooms used by any algorithm (including greedy) $\geq \text{depth}$, so greedy uses exactly as many classrooms as the depth.

Interval Graphs

NOT IN SYLLABUS

- Interval scheduling and interval partitioning can be seen as graph problems
- **Input**
 - Graph $G = (V, E)$
 - Vertices V = jobs/lectures
 - Edge $(i, j) \in E$ if jobs i and j are incompatible
- Interval scheduling = **maximum independent set (MIS)**
- Interval partitioning = **graph coloring**

Interval Graphs

NOT IN SYLLABUS

- MIS and graph coloring are NP-hard for general graphs
- But they're efficiently solvable for “**interval graphs**”
 - Graphs which can be obtained from incompatibility of intervals
 - In fact, this holds even when we are not given an interval representation of the graph
- Can we extend this result further?
 - Yes! Chordal graphs
 - Every cycle with 4 or more vertices has a chord

